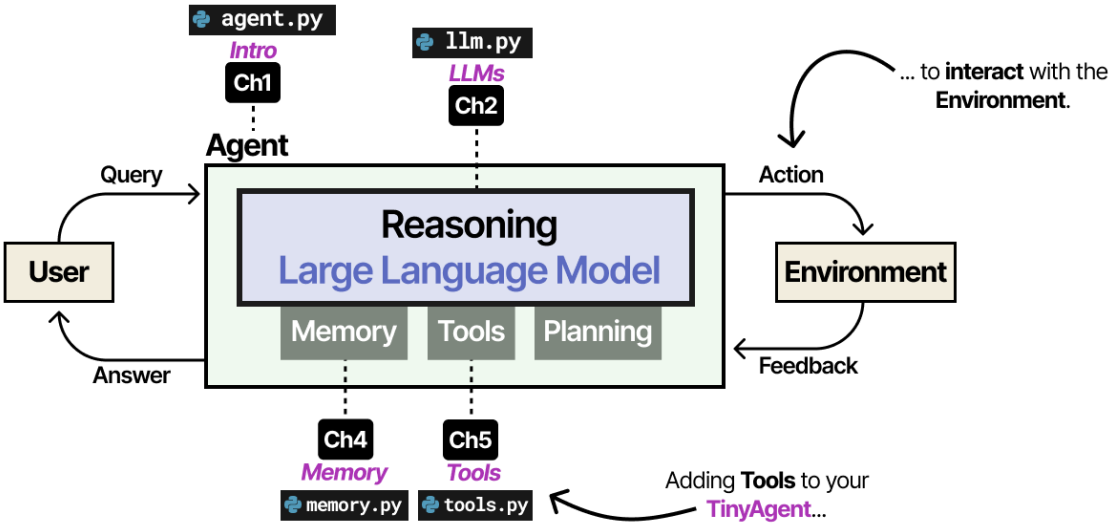
Table of Contents

Module 4 — Tools & MCP (Lecture Notes)

Course: 307307 BI Methods for Generative AI — *Agentic AI half* | **University of Petra Reference:** Grootendorst & Alammar, *An Illustrated Guide to AI Agents* (O'Reilly, 2026), Chapter 5.

> So far our agent can think (Module 2) and remember (Module 3) — but it still can't *do* anything. This module gives it **tools**, so it can act on the world: use a calculator, search the web, read a file. In the lab you give the `TinyAgent` a tool and watch it actually call it.



Adding Tools to the TinyAgent. An Illustrated Guide to AI Agents (O'Reilly, 2026).

Learning Objectives

By the end of this module you should be able to:

- Explain what a **tool** is and why an LLM needs one to *act*.
- Trace the **tool-calling flow**: create → define → select → call → process output.
- Explain **intent vs action** — and why executing the tool ourselves is a **safety boundary**.
- Describe how a tool is **defined** to the LLM (prompted vs native / schemas).
- Explain what **MCP** is and the problem it solves; recognise **Skills**.

1. What is a tool, and why?

A plain LLM is a text-in / text-out function. It can *describe* wanting to search the web or do a calculation, but it **cannot act**. A **tool** is simply a function or API we let the agent use — a calculator, a web search, a file reader, a database query. Tools fall into two kinds:

- **Take an action** (send an email, write a file), and

- **Retrieve information** (search, read a document, query an API).

Without tools, an agent can only think and plan — not act.

2. The tool-calling flow

Using a tool is more than "the LLM runs it." There are five steps:

1. **Create** the tool — write the function (well-documented; the docs are what the LLM reads).
2. **Define** it to the LLM — tell the model the tool exists, what it does, and how to call it.
3. **Select** — the LLM decides *which* tool (if any) fits the task (reasoning helps here).
4. **Call** — the LLM outputs a **tool call** (a structured string naming the tool and arguments).
5. **Process output** — *we* run the tool and feed the result back to the LLM.

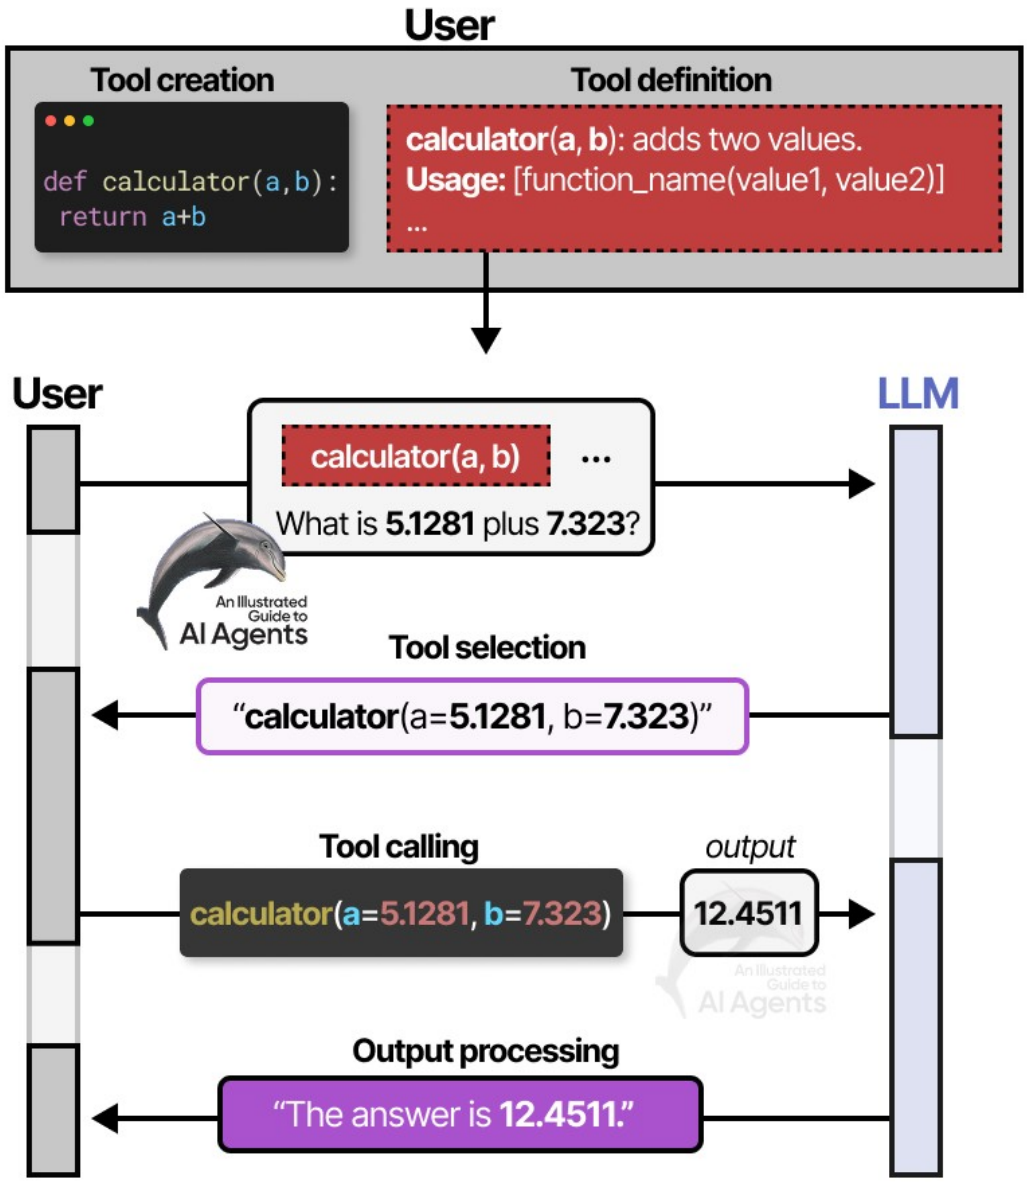


Figure 5-5 — the tool-calling flow. An Illustrated Guide to AI Agents (O'Reilly, 2026).

3. Intent vs action — the safety boundary

This is the key idea. When the LLM "calls" a tool, it only **outputs text** expressing the *intent* — for example `{"tool": "multiply", "kwargs": {"a": "5.1", "b": "7.3"}}`. **Nothing happens** until *our* program reads that text, runs the real function, and returns the result.

That gap is a feature, not a bug:

- The model's request to act is **not** permission to act — our code decides whether to run it.
- We can require **human approval** for risky tools (delete a file, send a payment) before they run.

Analogy: the LLM is like an assistant who writes "please wire $500 to this account" on a sticky note. Nothing moves until a person (your code) actually does it — and can refuse.

4. How the LLM is told about tools

- **Prompting (in-context)** — list the tools in the **system prompt** and tell the model the format to reply in (e.g. "reply with JSON `{\"tool\": ..., \"kwargs\": ...}`"). Simple and transparent; this is what we use in the lab. The small piece of code that reads the reply and extracts the call is called a **parser**.
- **Native / structured schemas** — newer models are trained to accept a formal **JSON schema** of each tool (name, description, parameters) and emit tool calls directly. Cleaner, more reliable, but more of a "black box."

Either way, the LLM must still *choose* and *format* the call correctly — which is why reasoning models are better at it.

Good practice: document tools well; keep the number of tools small (fewer than ~10 is a good start); keep each tool's scope narrow (few parameters).

5. The THOUGHT → ACTION → OBSERVATION loop

After the tool runs, we add its result back to the conversation (as an **OBSERVATION**) and let the LLM continue. You'll see these words in the lab: **THOUGHT** (what the model decides), **ACTION** (the tool call), **OBSERVATION** (the tool's result), **ANSWER** (the final reply). Looping over these is the **ReAct** pattern — the heart of the next module (Planning).

6. MCP — the Model Context Protocol

If every model talks to tools differently and every tool exposes itself differently, you get the **N × M problem**: N models × M tools = a custom integration for every pair. **MCP** (an open standard from Anthropic, late 2024) fixes this by standardising the connection — often called the **"USB-C port of AI."** With MCP it becomes **N + M** connections, and the *tool provider* maintains the integration, not you.

MCP has three parts:

- **Host** — the app using the LLM (for us, the TinyAgent; in practice Cursor, Claude, Copilot).
- **Client** — lives in the host; keeps a one-to-one connection to a server.
- **Server** — exposes a tool/data source (GitHub, arXiv, your files) via the standard.

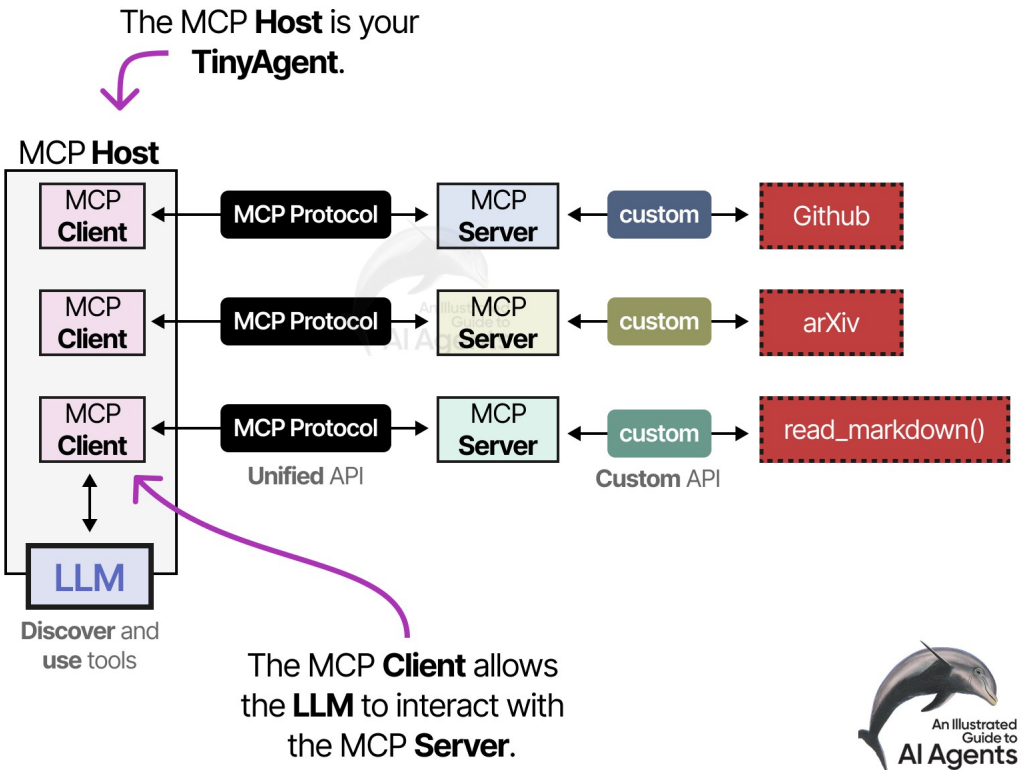


Figure 5-29 — the MCP host, client, and server. An Illustrated Guide to AI Agents (O'Reilly, 2026).

There are MCP servers for GitHub, Home Assistant, and many more. A related standard, **A2A (Agent2Agent)**, standardises communication *between agents*.

7. Skills (briefly)

Tools tell an agent *what it can do*; **Skills** tell it *how to get a recurring job done* — a recipe of steps, conventions, and references for your specific context (also an Anthropic idea). Skills use **progressive disclosure**: only the name + description are always loaded; the full instructions load when the agent decides the skill is relevant; extra files load only when needed. This keeps the context window lean. *(This is exactly the mechanism behind the skills you may already use with Claude.)*

8. Common Misconceptions

- **"The LLM runs the tool."** No — it only emits the *intent*; our code runs it. (That's the safety boundary.)
- **"More tools = more capable."** Too many tools make selection harder and the context heavier. Keep the set small.
- **"MCP is an AI model / makes the model smarter."** No — MCP is a *connection standard*; the model still needs tool-calling ability.
- **"A tool call is guaranteed correct."** The model can pick the wrong tool or mis-format the call; validate and handle failures.

9. Summary

- A **tool** lets the agent *act*; without tools it can only think.

- Tool use has five steps: create, define, select, call, process output.
- The LLM only emits the **intent**; *we* execute — a built-in **safety boundary** (and where human-in-the-loop lives).
- Tools are defined by **prompting** or **native schemas**.
- **MCP** standardises tool connections (the "USB-C of AI"); **Skills** package procedural know-how.

Review Questions

1. Why can't a plain LLM use a tool on its own? What does it actually produce?
2. List the five steps of the tool-calling flow.
3. Explain "intent vs action" and why it is a safety boundary. Where would human approval fit?
4. What problem does MCP solve, and what are its three components?
5. What is a Skill, and what does "progressive disclosure" mean?

Figures in the slides are reproduced from Grootendorst & Alammar, An Illustrated Guide to AI Agents (O'Reilly, 2026), for internal University of Petra classroom use.