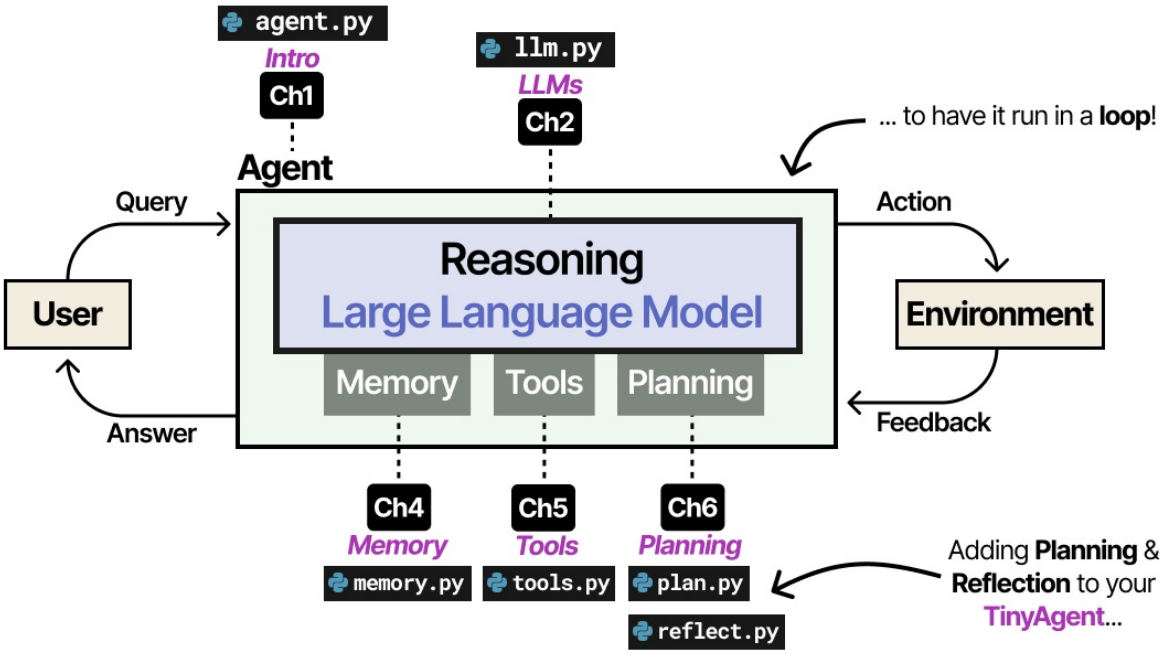
Module 5 — Planning & Reflection (ReAct) (Lecture Notes)

Course: 307307 BI Methods for Generative AI — *Agentic AI half* | **University of Petra Reference:** Grootendorst & Alammar, *An Illustrated Guide to AI Agents* (O'Reilly, 2026), Chapter 6 (agentic RAG: Chapter 4).

> Our agent can think (M2), remember (M3), and act (M4) — but only **one step at a time**. This module adds the final ingredient: **planning** and the **ReAct loop**, which let the agent break a task into steps and work through them **on its own**. This is what finally makes it an *agent*. In the lab you watch it solve a multi-step problem autonomously.



Adding Planning & Reflection to the TinyAgent — now it runs in a loop. An Illustrated Guide to AI Agents (O'Reilly, 2026).

Learning Objectives

By the end of this module you should be able to:

- Explain **task decomposition** — breaking a goal into smaller steps.
- Describe the **ReAct** loop: THOUGHT → ACTION → OBSERVATION, repeated until done.
- Explain what makes an agent **autonomous** (a loop + a stopping rule).
- Describe **reflection** (self-critique) and why it helps.
- Explain **agentic RAG**: retrieval as a *tool* inside the ReAct loop, and how it differs from the fixed RAG pipeline of Module 3.
- Recognise what lies **beyond a single agent**: evaluation, multi-agent, multimodal, coding agents.

1. Planning — decompose the task

Real tasks take many steps. Ask an agent to "add a feature to my code" and it must clarify requirements, read the codebase, design, implement, test, and document. **Task decomposition** is breaking the big goal into these smaller, ordered sub-tasks so the agent knows what to do next — and can **revise the plan** as it learns more.

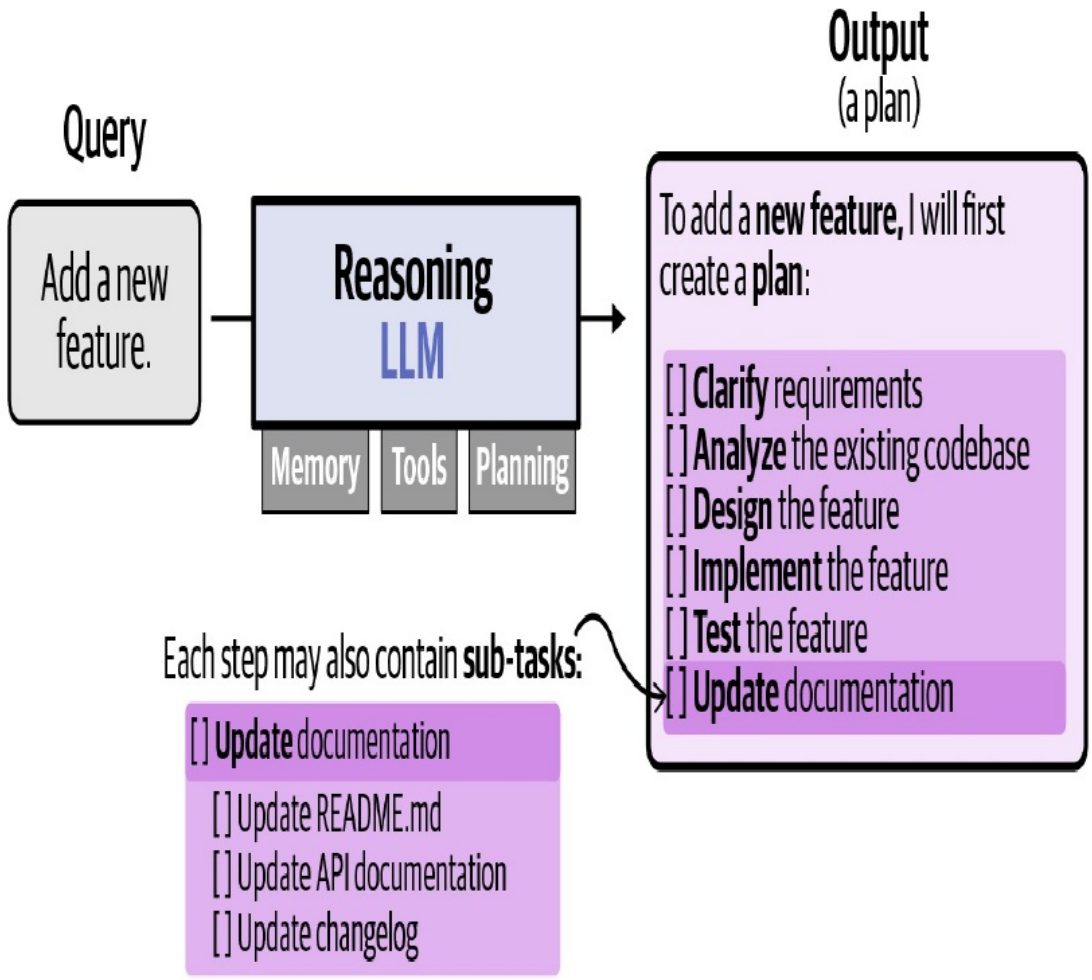


Figure 6-2 — the LLM turns a query into a plan. An Illustrated Guide to AI Agents (O'Reilly, 2026).

A reasoning LLM does this naturally (it "thinks step by step"). Prompting tricks like "*Let's think step by step*" (chain-of-thought) or "*make a plan, then solve it*" (plan-and-solve) encourage it.

2. ReAct — Reason + Act in a loop

Planning alone isn't enough: the agent must **act**, see what happened, and decide the next move. **ReAct** (Reason and Act) interleaves thinking and acting in a repeating cycle:

- **THOUGHT** — the agent reasons about what to do next.
- **ACTION** — it calls a tool (from Module 4).
- **OBSERVATION** — we run the tool and feed the result back.

...and it repeats, using each observation to decide the next thought — until it produces a final answer.

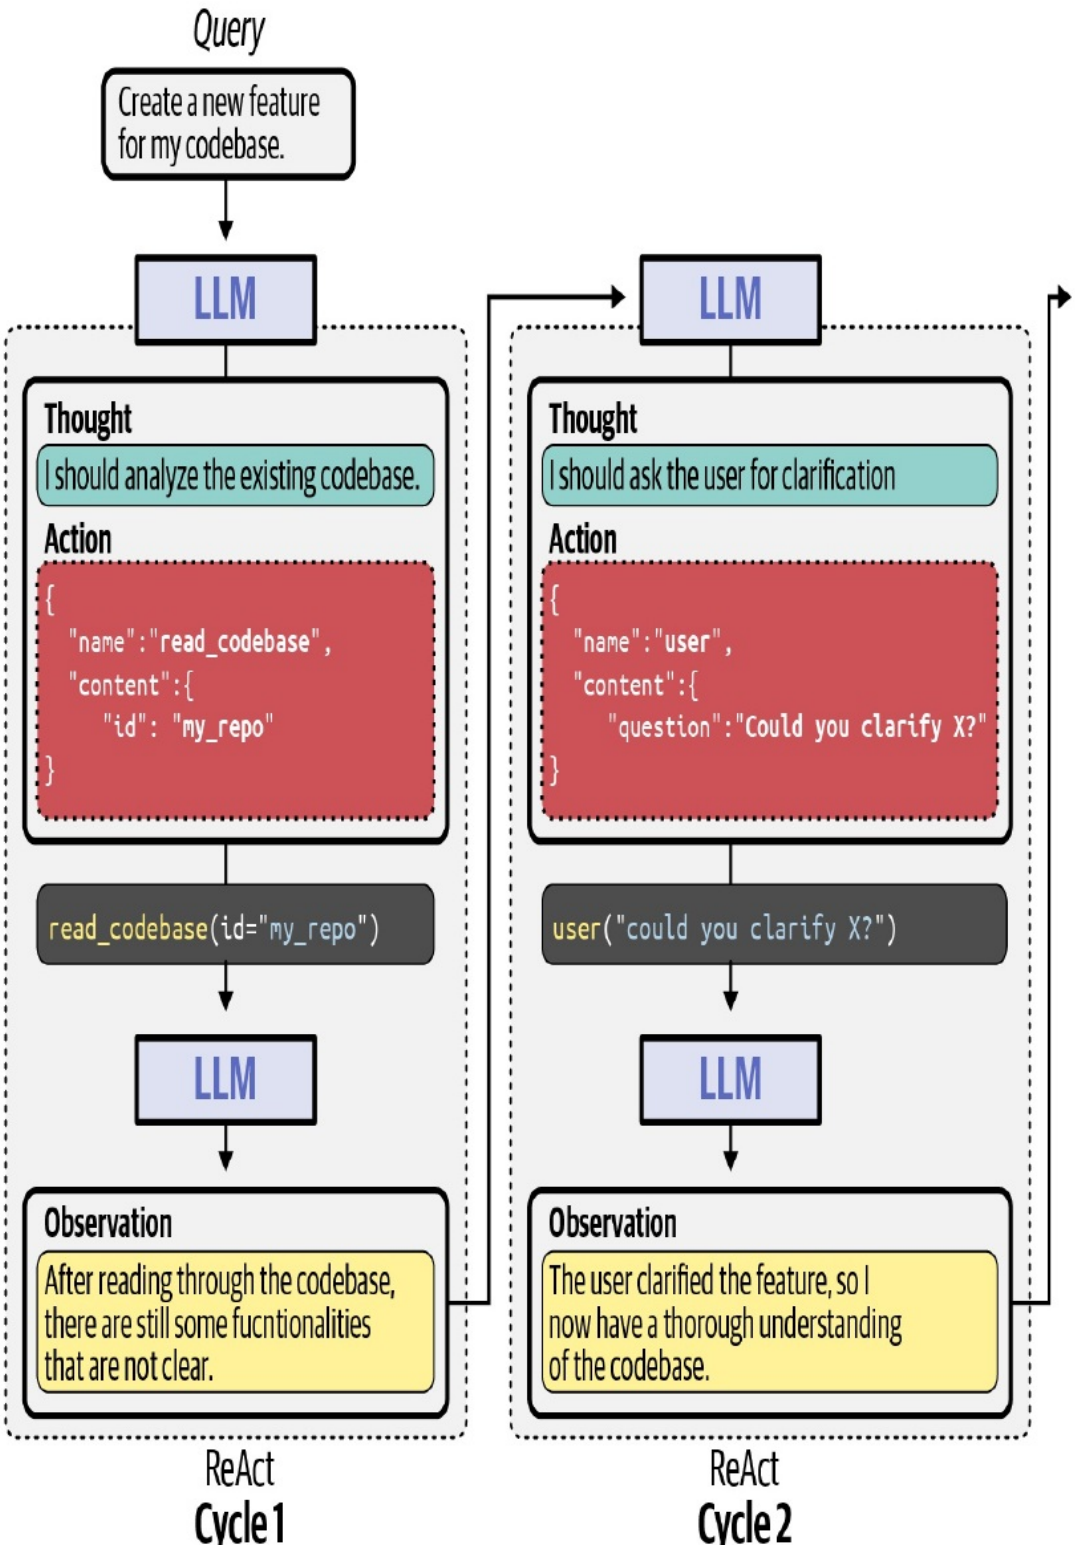


Figure 6-13 — two ReAct cycles. An Illustrated Guide to AI Agents (O'Reilly, 2026).

3. What makes it autonomous

The "magic" is surprisingly simple — **a loop**. The agent keeps running THOUGHT → ACTION → OBSERVATION until it decides to stop (by calling a special `final_answer`). Two safeguards matter:

- **A stopping rule** — the agent signals it's done with `final_answer`.
- **`max_steps`** — a cap on how many loops it may run, so a confused agent can't loop forever.

Put together, the four pieces from the whole course now form a real agent:

- **Reasoning LLM** (the brain) · **Memory** (so it remembers its steps) · **Tools** (so it can act) · **Planning / ReAct** (so it sequences actions autonomously).

That is the complete single-agent foundation — everything since Module 1 has been building to this.

4. Reflection — learning from its own mistakes

Agents can get stuck or go wrong. **Reflection** has the agent **critique its own work** and course-correct:

- **Self-Refine** — the same LLM drafts an answer, gives itself feedback, and refines it, in a loop (like a writer editing their own draft).
- **Reflexion** — a richer version with separate roles: one acts, one evaluates how well it did, one writes reflective feedback for the next attempt.

Reflection keeps an agent from getting stuck repeating the same mistake. (These are prompting techniques — straightforward to add.)

5. Agentic RAG — retrieval as a tool

In Module 3, **RAG** was a fixed pipeline: embed the question → retrieve **once** → add the results to the prompt → generate. The *developer's code* decides that retrieval happens, and the search query is simply the user's question. (The book calls this *vanilla* or *naive* RAG.)

That breaks down when **what to search for next depends on what the first search found**. Take a question about a fictional café: "*When can I find the store manager?*" The question contains no name. One search can find *who* the manager is, but only a **second** search, for that person's working hours, can answer the question. A one-shot pipeline cannot know to run it.

Agentic RAG makes retrieval a **tool** inside the ReAct loop, next to `add` or `multiply`. The agent decides:

- **what** to search for (it writes the query itself),
- **whether** to search again, using what it has just read (the OBSERVATION),
- **when** it has enough to give a `final_answer`.

	RAG (Module 3)	Agentic RAG
Who decides to retrieve?	The developer's code	The agent, at each step
Number of retrievals	Always one	As many as the agent needs (capped by `max_steps`)
Search query	The user's question	Written by the agent, can use earlier observations
Cost and speed	One LLM call	Several LLM calls: slower and more expensive
Predictability	High	Lower; read the trajectory

Nothing about the ReAct loop changes. The agent is given one more tool (`search_docs`) and a system prompt that describes it.

More than one knowledge source. Each source (a handbook, a price list, a web search) can be its own tool, and the agent chooses which to query: it acts as a *router*. In **multi-agent RAG**, smaller retrieval agents, each specialised in one source, are coordinated by a supervisor agent. The book weighs the trade-off (its Table 4-1): a single agent is cheaper and easier to debug but is a single point of failure; several agents are more modular and can check each other, at a higher cost and complexity. Its advice is to **start with a single agent** as a baseline. In this course multi-agent systems are only surveyed.

What agentic RAG does *not* fix. It adds facts to the prompt; it does not change the model. A model can still write a poor query, stop searching too early, or answer from its own memory instead of from what it retrieved. The retrieval itself can also fail, because a vector search returns its nearest documents even when none is relevant. That is why we keep `max_steps`, tell the agent to answer *only* from the search results, and read the printed THOUGHT / ACTION / OBSERVATION steps.

6. Beyond a single agent (course survey)

You've now built a complete **single** agent. The field goes further — here's the map of what comes next, at a glance:

- **Evaluation** — agents are hard to judge: look at the **outcome** (did the task get done?) *and* the **trajectory** (were the steps sound?), plus reliability and safety.
- **Multi-agent systems** — several specialised agents collaborate, often coordinated by a **supervisor** agent.
- **Multimodal agents** — agents that can also **see and hear** (images, audio, video), not just read text.
- **Coding agents** — agents that **run, test, and fix** real code in a dedicated environment (e.g. the assistants many developers now use).

These are surveyed briefly to show where the field is heading; the single-agent skills you built transfer directly.

7. Common Misconceptions

- **"ReAct is a special model."** No — it's a **loop + a prompt format**. Any capable LLM can do it.
- **"Planning means one fixed plan."** No — good agents **revise** the plan as they learn (reflection).
- **"An agent will always finish."** Not guaranteed — that's why we set `max_steps` and watch the trajectory.
- **"More steps = better."** Not necessarily; extra loops cost tokens and can wander. Stop when the goal is met.
- **"Agentic RAG is always better than RAG."** No. If the question is itself a good search query and the answer sits in one or two documents, plain RAG is cheaper, faster and more predictable. Use the agent when later searches depend on earlier results, or when there are several sources to choose from.
- **"The agent will always search well."** No. It can write a poor query, stop too early, or ignore the results. Check the trajectory.

8. Summary

- **Planning** decomposes a goal into ordered sub-tasks.
- **ReAct** loops THOUGHT → ACTION → OBSERVATION until a `final_answer`.
- A **loop + a stopping rule (`max_steps`)** is what makes the agent autonomous.
- **Reflection** lets an agent critique and improve its own work.
- **Agentic RAG** = retrieval as a tool in the ReAct loop: the agent decides what to search for and whether to search again, at the price of more LLM calls.
- Reasoning + Memory + Tools + Planning = a **complete single agent**. Beyond it: evaluation, multi-agent, multimodal, coding agents.

Review Questions

1. What is task decomposition, and why does an agent need it?
2. Name the three parts of a ReAct cycle and what each does.
3. What two things make an agent autonomous — and what does `max_steps` protect against?
4. What is reflection? Briefly contrast Self-Refine and Reflexion.
5. Name the four components that together make a complete single agent.
6. How does agentic RAG differ from the RAG pipeline of Module 3? Give an example of a question that needs more than one search.
7. When is plain RAG enough, and what do you pay for when you use agentic RAG instead?

Figures in the slides are reproduced from Grootendorst & Alammar, An Illustrated Guide to AI Agents (O'Reilly, 2026), for internal University of Petra classroom use.